

# Machine Learning

Pawel Wocjan

University of Central Florida

Spring 2019

# Layers: the building blocks of deep learning

- ▶ The fundamental data structure in neural networks is the **layer**.
- ▶ A layer is a data-processing module that takes as input one or more tensors and that outputs one or more tensors.
- ▶ Some layers are stateless, but more frequently layers have a state: the layers **weights**, one or several tensors learned with stochastic gradient descent, which together contain the network's **knowledge**.

# Layers: the building blocks of deep learning

- ▶ Different layers are appropriate for different tensor formats and different types of data processing.
- ▶ Simple vector data, stored in 2D tensors of shape (samples, features) is often processed by **densely connected layers**, also called **fully connected layers** (the Dense class in Keras).
- ▶ Sequence data, stored in 3D tensors of shape (samples, timesteps, features), is typically processed by **recurrent** layers such as **long-short term memory (LSTM)** layer.
- ▶ Image data, stored in 4D tensors, is usually processed by 2D **convolutional** layers (Con2D).

<https://keras.io/layers/core/>

<https://keras.io/layers/convolutional/>

<https://keras.io/layers/recurrent/>

# Layers: the building blocks of deep learning

- ▶ You can think of layers as LEGO bricks of deep learning.
- ▶ Building deep-learning models in Keras is done by combining compatible layers to form useful data-processing pipelines.
- ▶ Layer compatibility means that every layer will only accept input tensors of a certain shape and will return output tensors of a certain shape.
- ▶ When using Keras, you don't have to worry about compatibility, because the layers you add to your model are dynamically built to match the shape of the incoming layer.

# Layers: the building blocks of deep learning

```
from keras import models
from keras import layers

network = models.Sequential()

network.add(layers.Dense(512,
                          activation='relu',
                          input_shape=(28 * 28,)))

# no need to specify input_shape for second layer
network.add(layers.Dense(10,
                          activation='softmax'))
```

The second layer didn't receive an input shape argument – instead, it automatically inferred its input shape as being the output shape of the first layer.

# Models: networks of layers

- ▶ A deep-learning model is a directed, acyclic graph of layers.
- ▶ The most common topology is a linear stack of layers, mapping a single input to a single output. These can be implemented using `models.Sequential()`.
- ▶ Initially, we will only work with linear stacks of layers.
- ▶ Later, we will also look at other network topologies such as two-branch networks, multi-head networks, and inception blocks.

<https://keras.io/getting-started/sequential-model-guide/>

# Models: networks of layers

- ▶ The topology of a network defines a **hypothesis space**.
- ▶ By choosing a network topology, you constrain your **space of possibilities** (hypothesis space) to a specific series of tensor operations, mapping input data to output data.
- ▶ You'll be then searching for a good set of values for the weight tensor involved in these tensor operations using stochastic a variant of gradient descent.
- ▶ Picking the right network architecture is more art than a science. We will study explicit principles for building neural networks and develop intuition as to what works or doesn't for specific problems.

# Loss functions & optimizers: keys to configuring the learning process

- ▶ Once the network architecture is defined, you still need to do two things:
  - ▶ **Loss function (objective function)**  
The quantity that will be minimized during training. It represents a measure of success for that task at hand.  
<https://keras.io/losses/>
  - ▶ **Optimizer** Determines how the network will be updated based on the loss function. Implements a specific variant of stochastic gradient descent (SGD).  
<https://keras.io/optimizers/>



# Loss functions & optimizers: keys to configuring the learning process

- ▶ Choosing the right objective function for the right problem is extremely important: your network will take any shortcut it can, to minimize the loss.
- ▶ Fortunately, there are simple guidelines you can use to choose the correct loss for common problems such as classification, regression, and sequence prediction.

Problem type	Last-layer activation	Loss function
Binary classification	sigmoid	binary_crossentropy
Multiclass, single-label classification	softmax	categorical_crossentropy
Multiclass, multi-label classification	sigmoid	binary_crossentropy
Regression to arbitrary values	None	mse
Regression to values in $[0, 1]$	sigmoid	mse or binary_crossentropy