

Softmax and categorical cross entropy loss

Pawel Wocjan

February 11, 2019

Abstract

We define softmax and categorical cross entropy for multiclass classification

1 Classification

Problem type	Last-layer activation	Loss function
Binary classification	sigmoid	binary_crossentropy
Multiclass, single-label classification	softmax	categorical_crossentropy
Multiclass, multi-label classification	sigmoid	binary_crossentropy

2 Softmax activation function

Let z_1, \dots, z_m be the weighted inputs of the m neurons of the last layer. Each of the neurons corresponds to one of the m classes of the multiclass classification problem at hand. We consider the single-label situation. We convert the weighted input vector $\mathbf{z} = (z_1, \dots, z_m)^T \in \mathbb{R}^m$ into a probability vector $\mathbf{a} = (a_1, \dots, a_m)^T \in \mathbb{R}^n$ by applying the so-called softmax activation function.

For $k \in [m]$, the activation a_k of the k th neuron is defined as follows:

$$a_k = \frac{e^{z_k}}{\sum_{j=1}^n e^{z_j}}. \quad (1)$$

It is straightforward to verify that this yields a probability distribution. The values a_k are all positive because the range of the exponential function is $(0, \infty)$. They sum up to 1 because of the normalization in the denominator.

Using the product and chain rule, we can show that

$$\frac{\partial a_k}{\partial z_j} = a_k \cdot (\delta_{jk} - a_j), \quad (2)$$

where δ_{jk} is the so called Kronecker delta, which is equal to 1 if $j = k$ and 0 if $j \neq k$.

3 Cross entropy

Let $\mathbf{p} = (p_1, \dots, p_m)$ and $\mathbf{q} = (q_1, \dots, q_m)$ be two probability distributions. In information theory, the cross entropy is defined by

$$H(\mathbf{p}, \mathbf{q}) = - \sum_{k=1}^m p_k \log q_k. \quad (3)$$

See https://en.wikipedia.org/wiki/Cross_entropy for a quick overview.

4 Categorical cross entropy loss function

Let $y \in [m]$ be a label. Using the so-called one-hot or categorical encoding, we construct a corresponding vector $\mathbf{y} = (y_1, \dots, y_m)^T \in \mathbb{R}^m$ such that $m - 1$ of its entries are equal to 0 and exactly one entry is equal to 1. The position of the entry 1 is given by the label y . For instance, for $m = 3$ classes, label 1 corresponds to $(1, 0, 0)$, the label 2 to $(0, 1, 0)$, and the label 3 to $(0, 0, 1)$.

Assume that the feature vector \mathbf{x} produces the activation vector $\mathbf{a} = (a_1, \dots, a_m)$ in the last layer. Assume that the correct label is y . Then the categorical cross entropy loss \mathcal{L} is defined by

$$\mathcal{L} = - \sum_{k=1}^m y_k \log a_k, \quad (4)$$

where $\mathbf{y} = (y_1, \dots, y_m)$ is the categorical encoding of y . Observe that $\mathcal{L} = H(\mathbf{y}, \mathbf{a})$.

The partial derivatives of \mathcal{L} with respect to a_k is

$$\frac{\partial \mathcal{L}}{\partial a_k} = - \frac{y_k}{a_k}. \quad (5)$$

The partial derivatives of \mathcal{L} with respect to z_j is

$$\frac{\partial \mathcal{L}}{\partial z_j} = \sum_{k=1}^m \frac{\partial \mathcal{L}}{\partial a_k} \cdot \frac{\partial a_k}{\partial z_j} \quad (6)$$

$$= \sum_{k=1}^m - \frac{y_k}{a_k} \cdot a_k \cdot (\delta_{kj} - a_j) \quad (7)$$

$$= \sum_{k=1}^m y_k \cdot (a_j - \delta_{kj}) \quad (8)$$

5 Simple neural network with softmax activation and categorical cross entropy

Let us consider a network that takes feature vectors of the form $\mathbf{x} = (x_1, \dots, x_n)^T \in \mathbb{R}^n$ as input and has m output neurons with softmax activation. Then for $j \in [m]$, the weighted

input of the j th neuron is given by

$$z_j = \sum_{i=1}^n w_{ji}x_i + b_j. \quad (9)$$

The weights and the bias of j th neuron are (w_{j1}, \dots, w_{jn}) and b_j , respectively. The weighted inputs z_j are transformed into probabilities by the softmax activation.

Compute the partial derivatives of the categorical cross entropy loss function the weights and biases. You will need these to implement stochastic gradient descent for the homework. This should be very easy. The more difficult part was computing the partial derivatives of the categorical entropy loss function with respect to the weighted inputs.